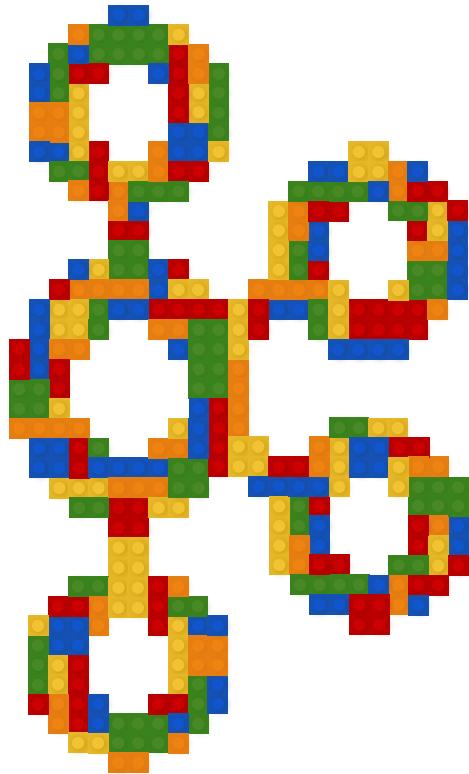


Brick-by-Brick

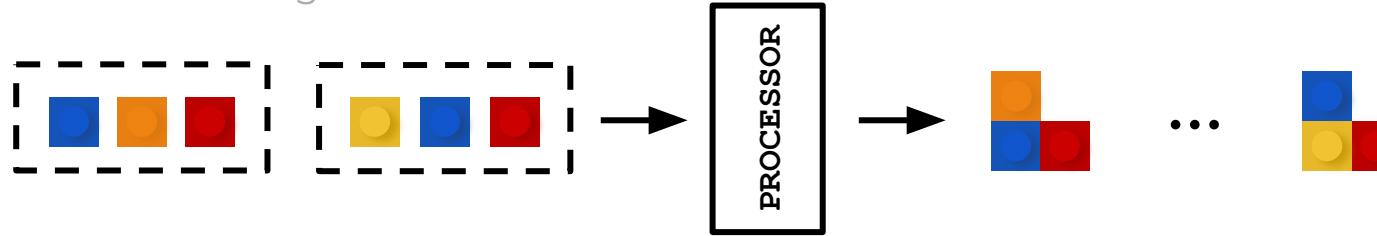
Exploring the Elements of Apache Kafka®



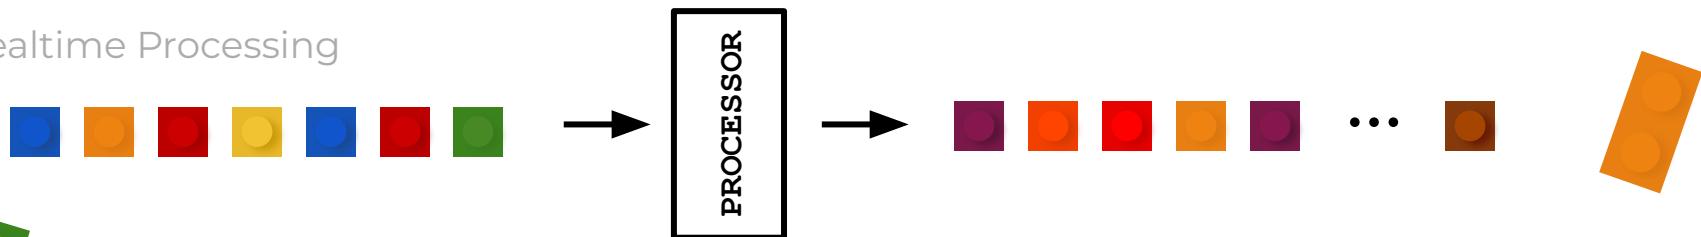
Why
Kafka?

... because real-time *makes sense*.

Batch Processing

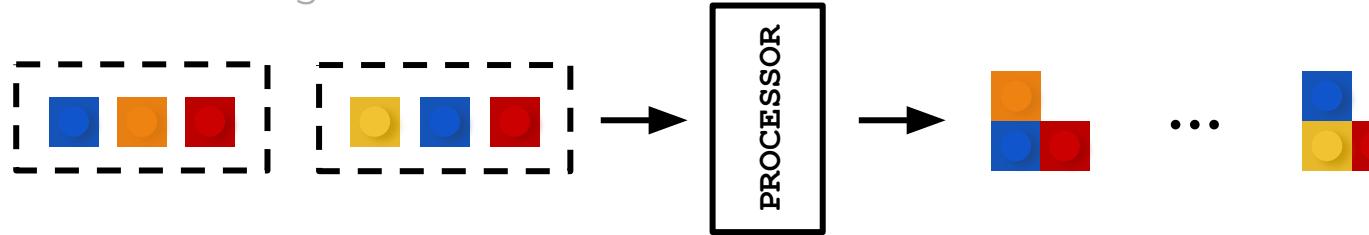


Realtime Processing

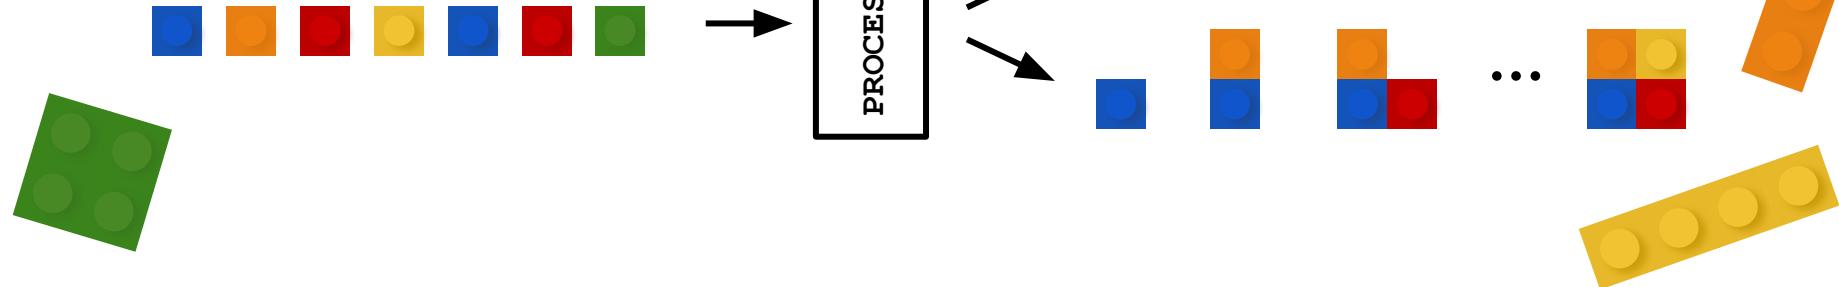


... because real-time *makes sense*.

Batch Processing

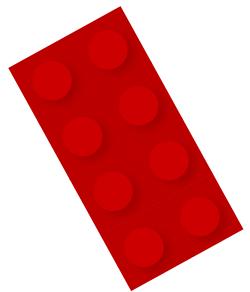


Realtime Processing

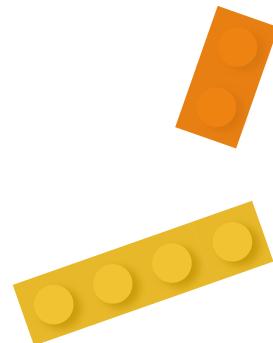


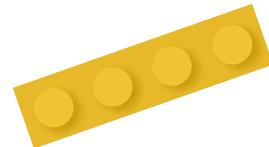
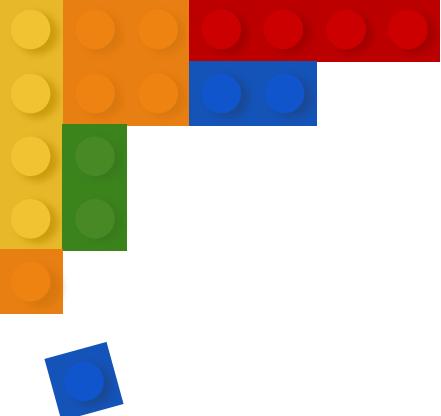
Kafka Components

Events



- Natural way to reason about things
- Indicate that something *has happened*
 - When
 - What/Who
- event = notification + state



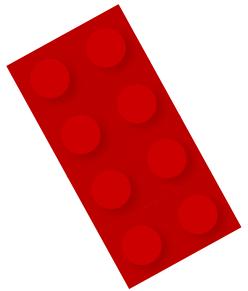


dfine@confluent.io

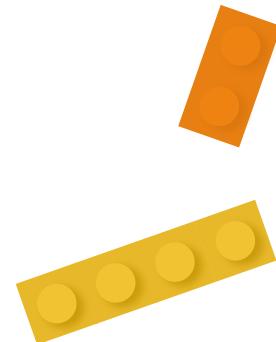
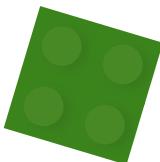
@TheDanicaFine

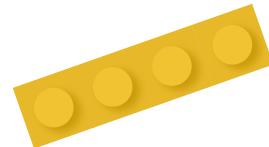
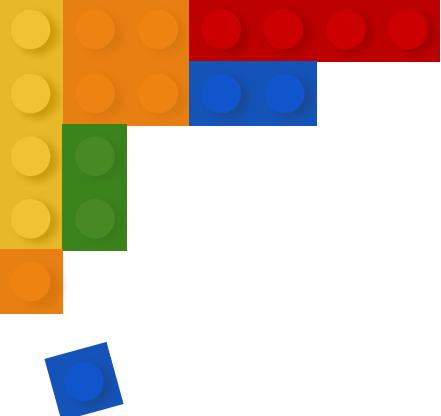
linkedin.com/in/danica-fine/

Events

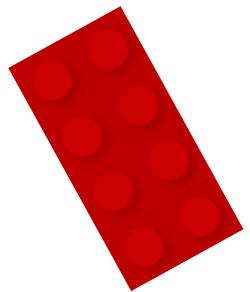


- Natural way to reason about things
- Indicate that something *has happened*
 - When
 - What/Who
- event = notification + state
- Immutable pieces of information

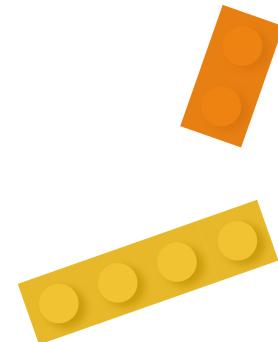




Events



- Natural way to reason about things
- Indicate that something *has happened*
 - When
 - What/Who
- event = notification + state
- Immutable pieces of information





Messages and Records

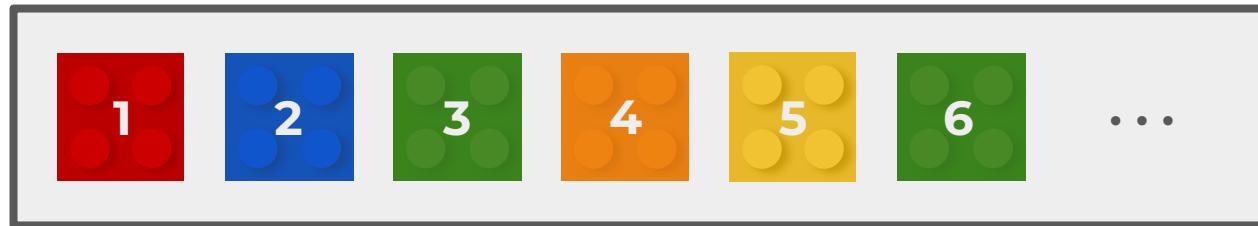
An aside.

Kafka Storage

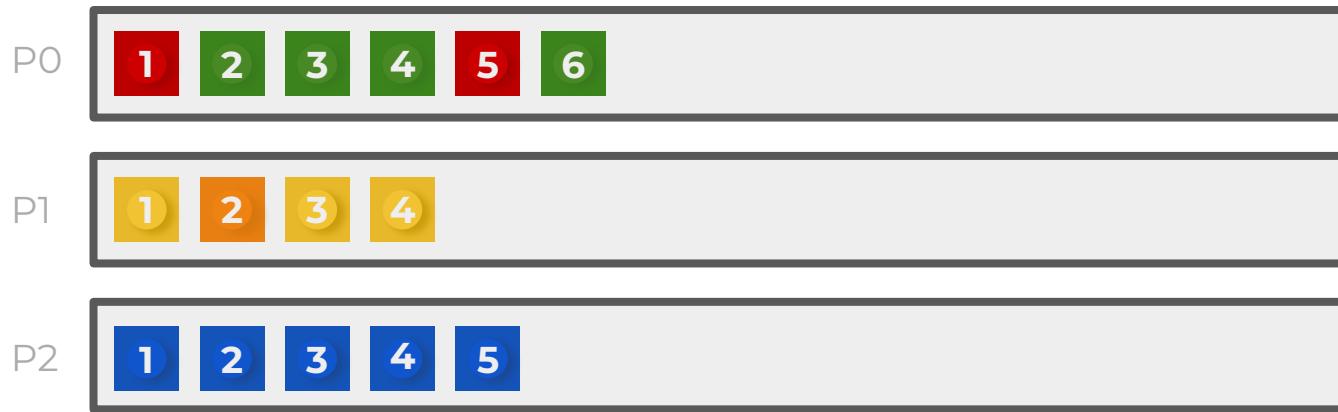


The Topic

A durable, immutable, append-only log of events.

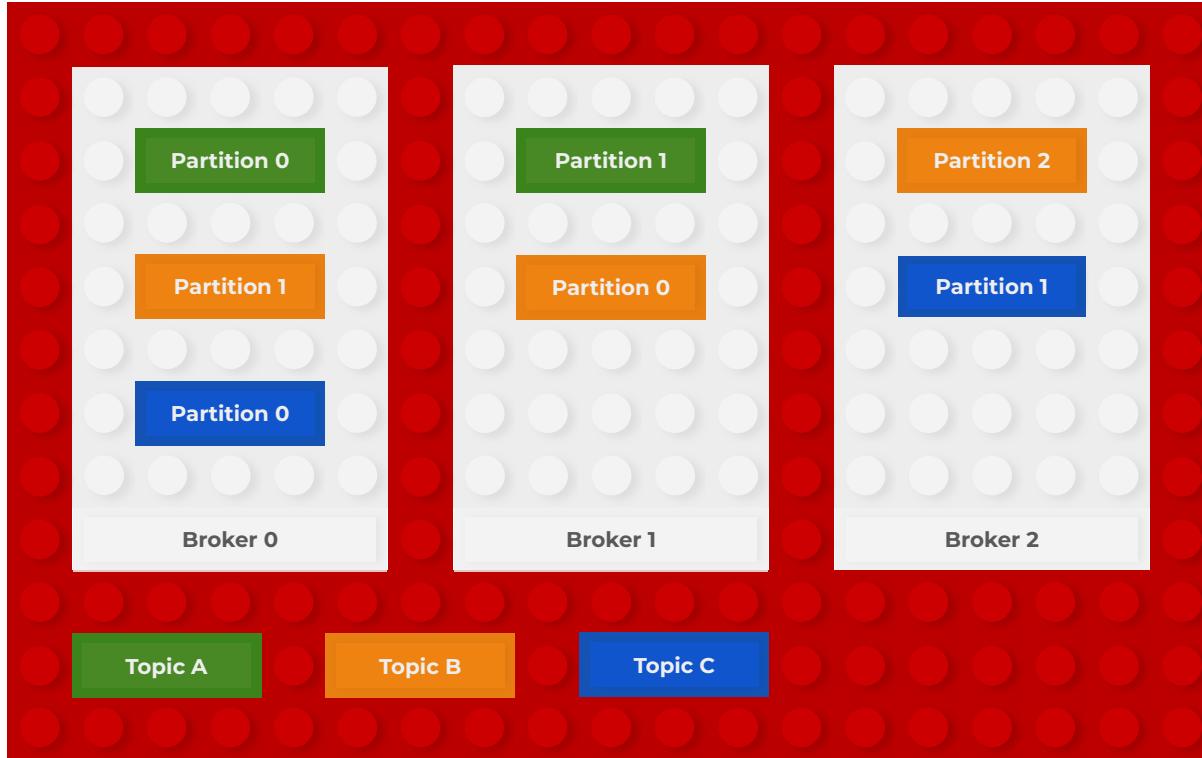


Topic-Partitions

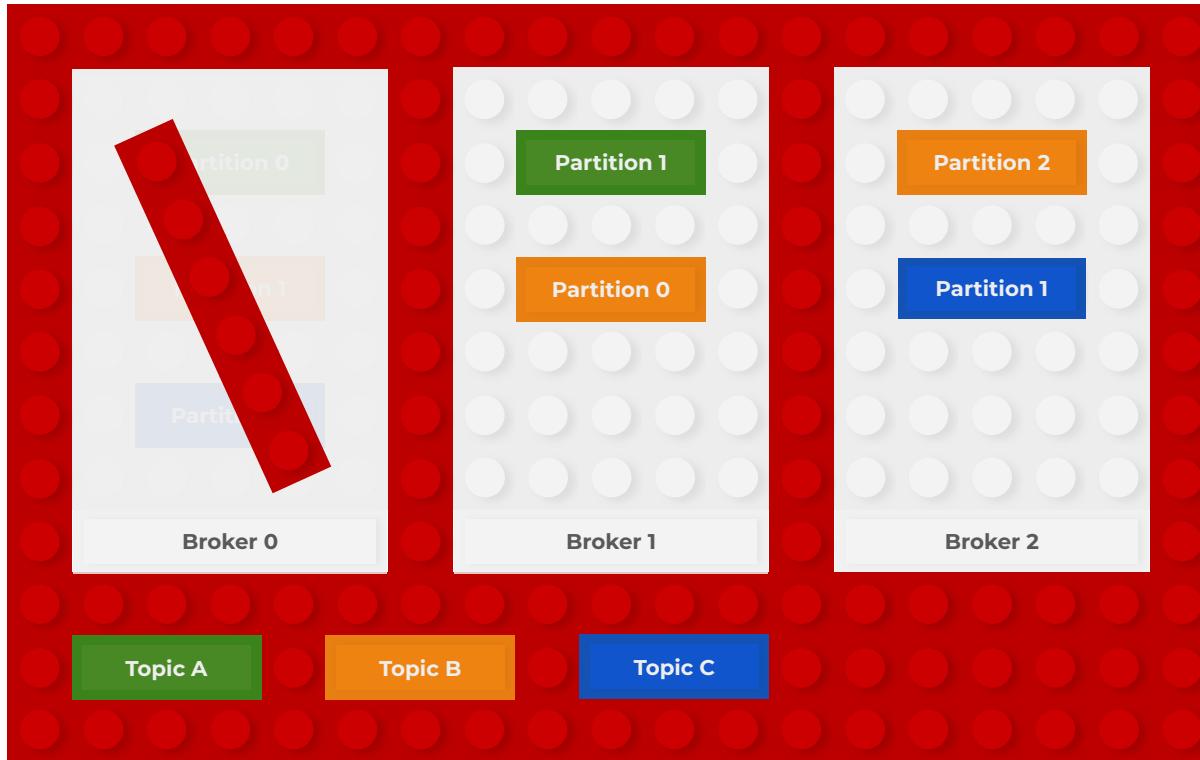


Why partitions?

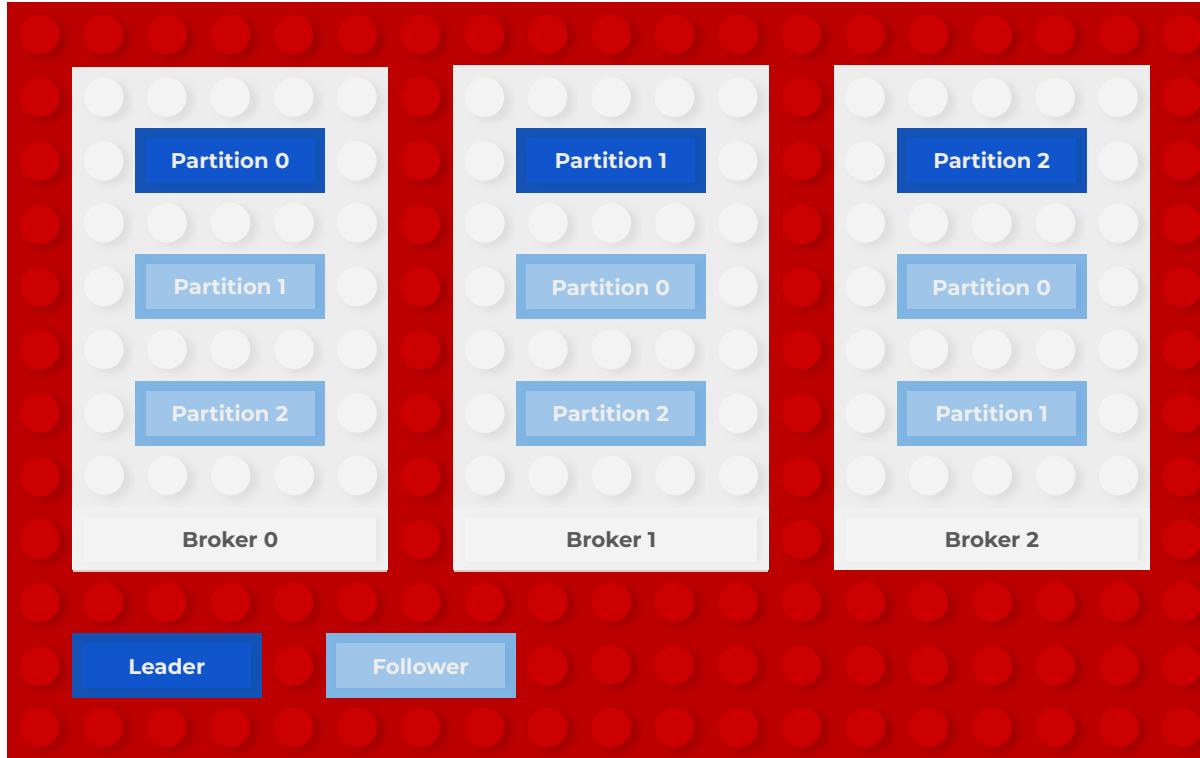
A Kafka Cluster



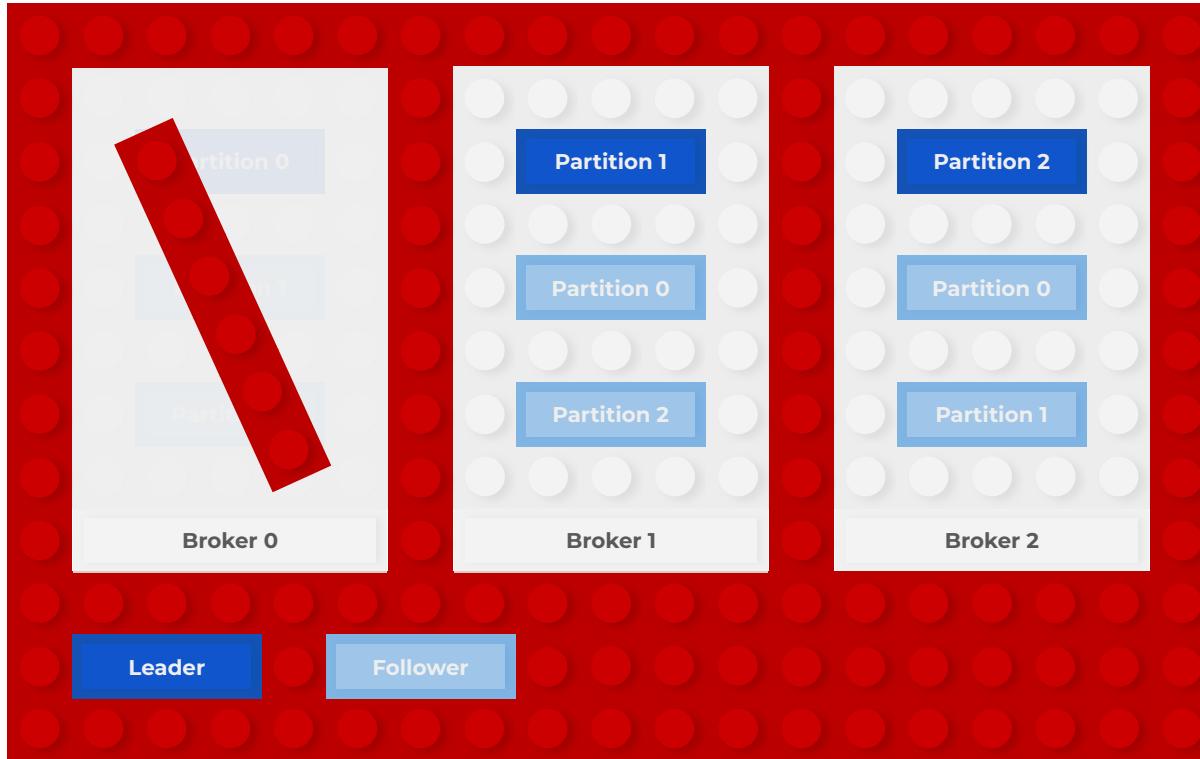
A Kafka Cluster



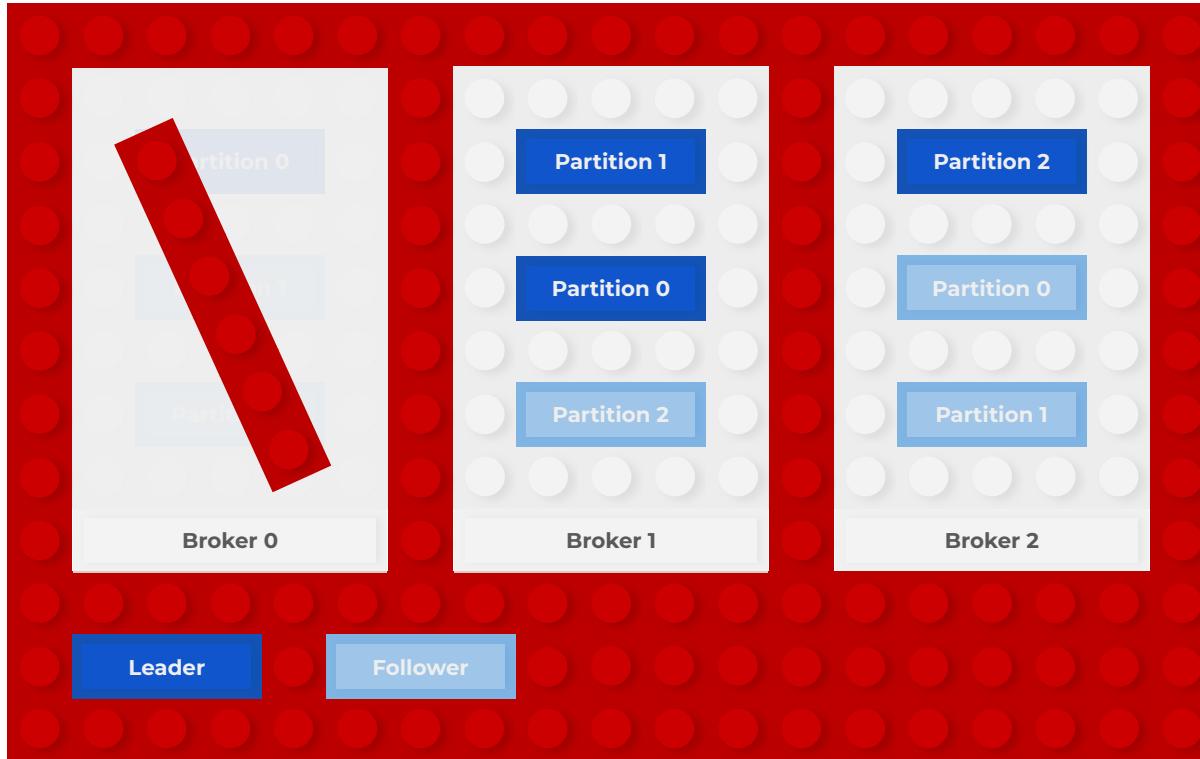
Replication



Replication



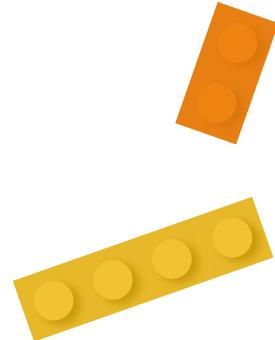
Replication





A record

- Required Fields
 - Topic
 - Value
- [Optional] Fields
 - Partition
 - Timestamp
 - Key
 - Headers

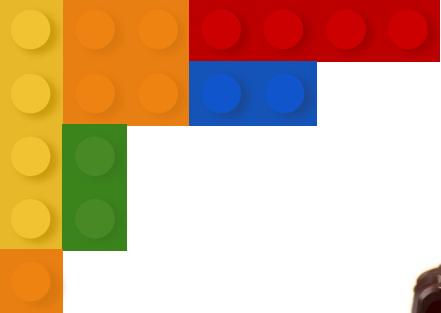




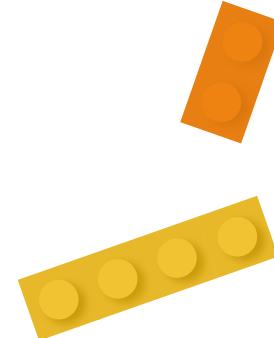
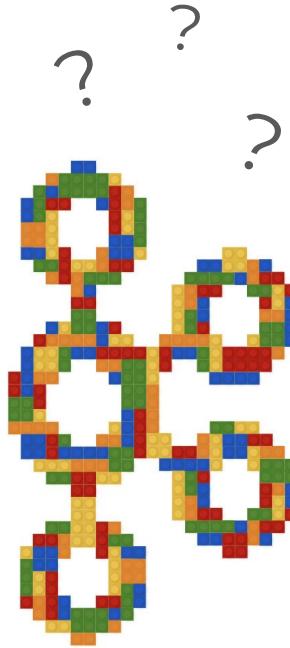
A record

- Required Fields
 - Topic: 'cart-updates'
 - Value: 'Add 1 Blue 1x1 Plate to Cart 1234'
- [Optional] Fields
 - Partition: None
 - Timestamp: None
 - Key: 'Cart 1234'
 - Headers: None



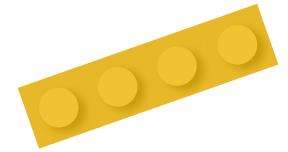


Producer





Brokers and Bytes

- Store bytes is more efficient
 - Serialization formats:
 - Avro
 - JSON
 - Protobuf
 - Configure producers to serialize the object
- 
- 

Schemas

```
{  
  "doc": "Pick a Brick cart update",  
  "fields": [  
    {  
      "name": "cart_id",  
      "type": "int"  
    },  
    {  
      "name": "cart_action",  
      "type": {  
        "name": "Status",  
        "type": "cart_action",  
        "symbols": ["ADD", "REMOVE",  
                   "UPDATE_QTY"]  
      }  
    },  
  ]  
}
```

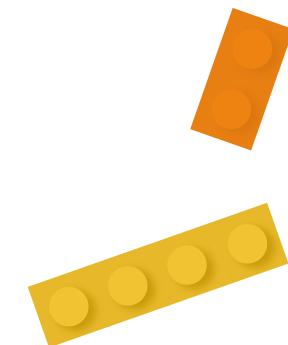
```
{  
  "name": "element_id",  
  "type": "int"  
},  
{  
  "name": "element_desc",  
  "type": "string"  
},  
{  
  "name": "quantity",  
  "type": "int"  
}  
,  
{  
  "name": "cart_updates",  
  "namespace": "lego",  
  "type": "record"  
}
```



Producing a Record: A recap



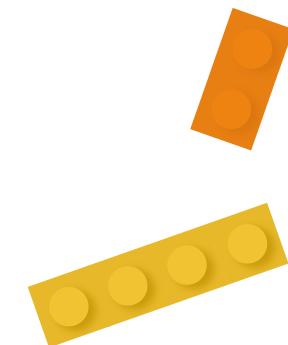
```
key: 1234  
  
{  
  "cart_id": 1234,  
  "cart_action": "ADD",  
  "element_id": 302423,  
  "element_desc": "Bright Blue  
    1x1 Plate",  
  "quantity": 1  
}
```

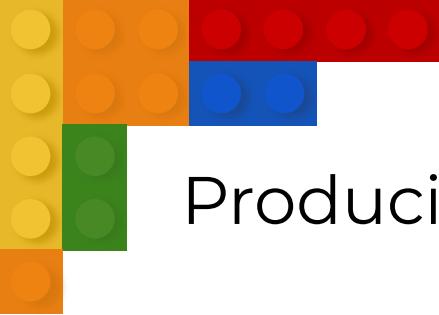




Producing a Record: A recap



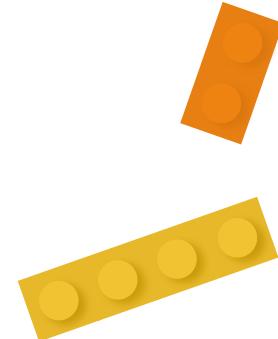
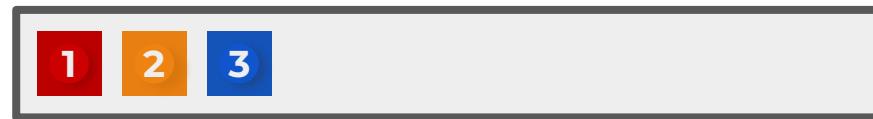
1. Serialize
 2. Partition
 3. Batch
 4. [Compress]
 5. Request
- 



Producing a Record: A recap



Cart Updates Topic

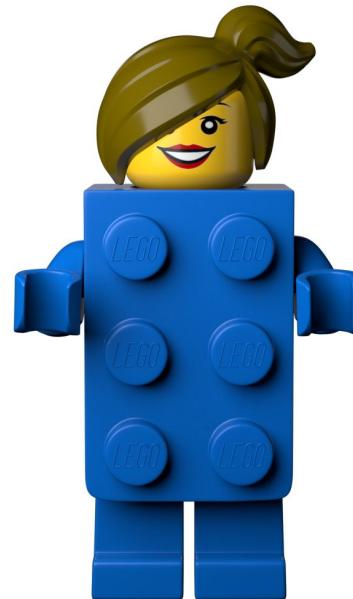




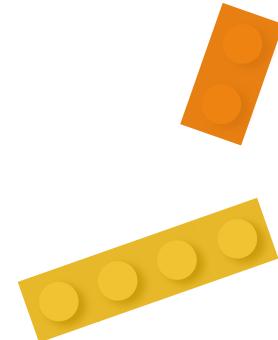
Making it More Interesting



Producer

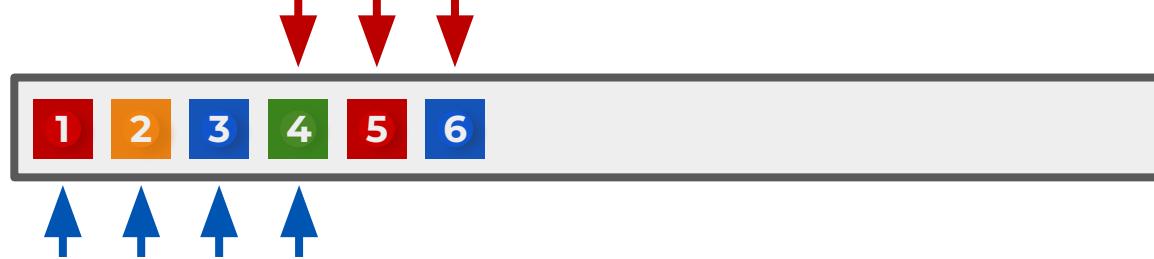


Consumer





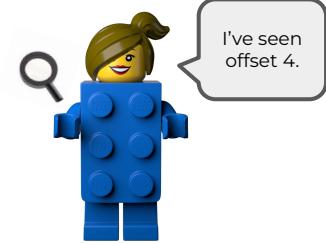
Producer



I've seen offset 4.



Consumer



Consumer 1



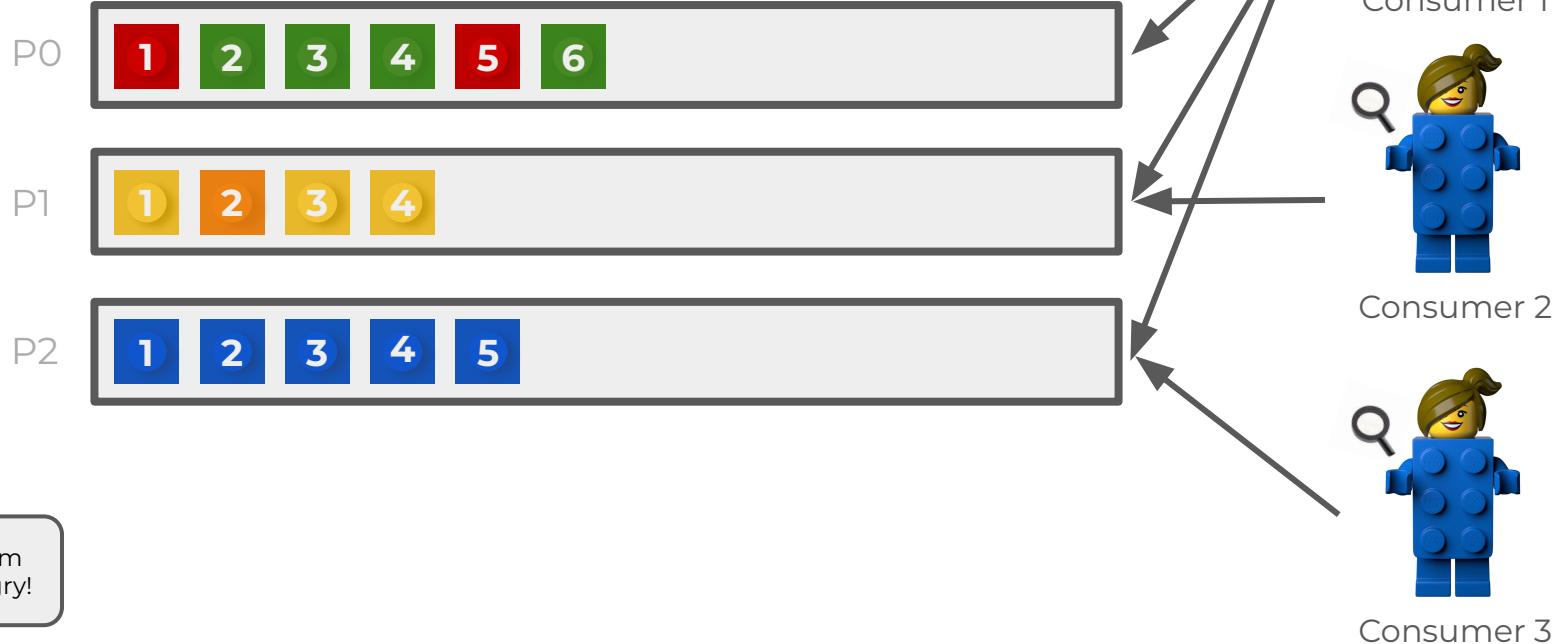
Consumer 2



Consumer 3



A Consumer Group





Break Time

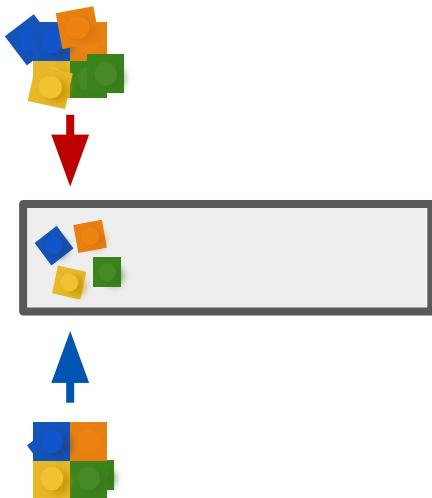
And a Summary

Kafka Ecosystem

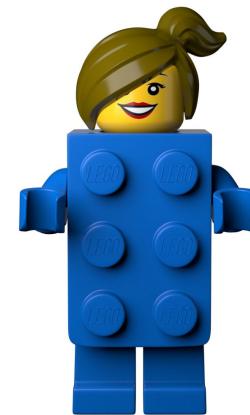
```
{
  "doc": "Pick a Brick cart update",
  "fields": [
    {
      "name": "cart_id",
      "type": "int"
    },
    ...
    {
      "name": "quantity",
      "type": "int"
    }
  ],
  "name": "cart_updates",
  "namespace": "com.picksbrick",
  "type": "record"
}
```



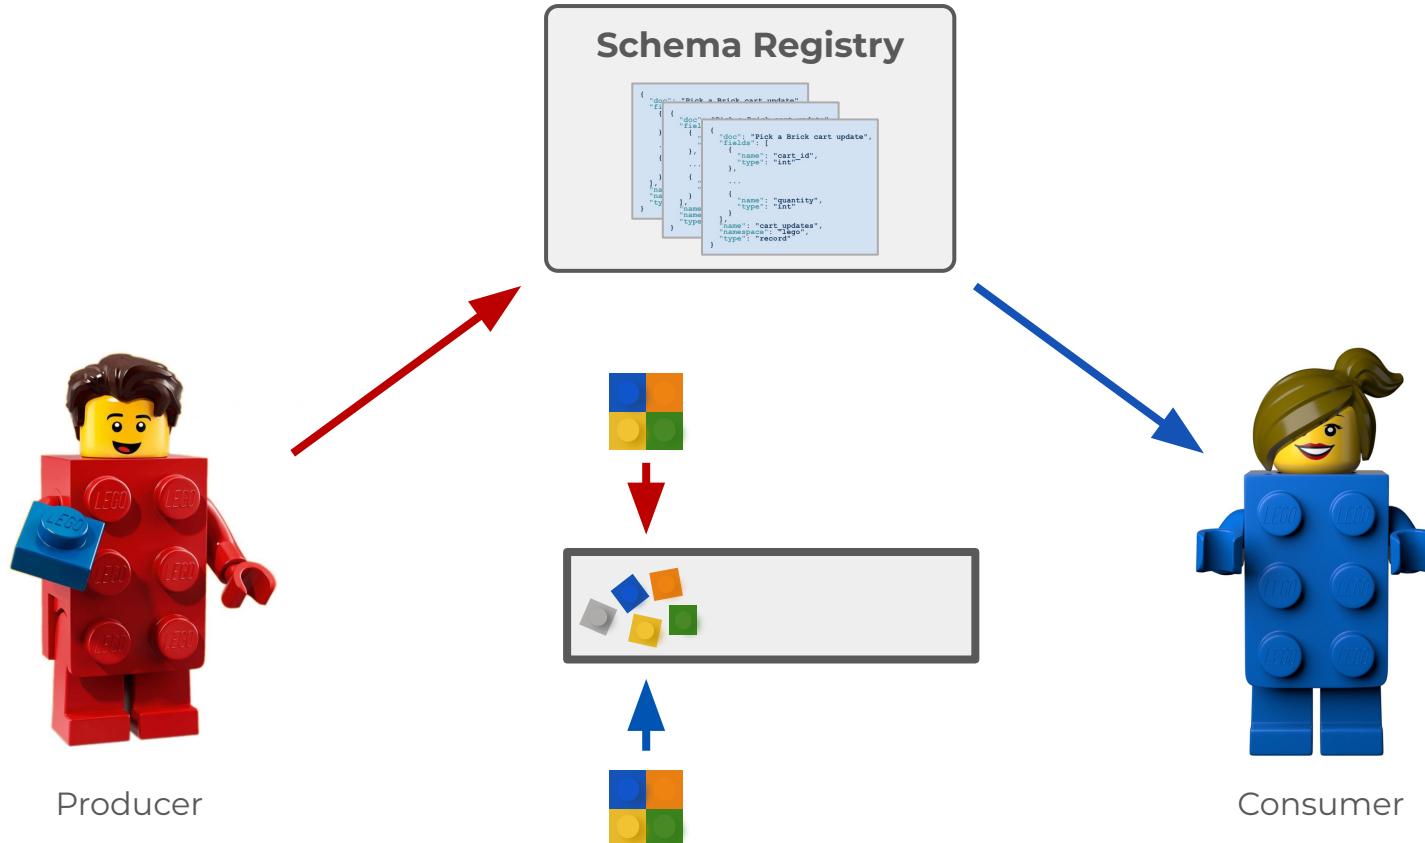
Producer



```
{
  "doc": "Pick a Brick cart update",
  "fields": [
    {
      "name": "cart_id",
      "type": "int"
    },
    ...
    {
      "name": "quantity",
      "type": "int"
    }
  ],
  "name": "cart_updates",
  "namespace": "com.picksbrick",
  "type": "record"
}
```



Consumer



Kafka Connect

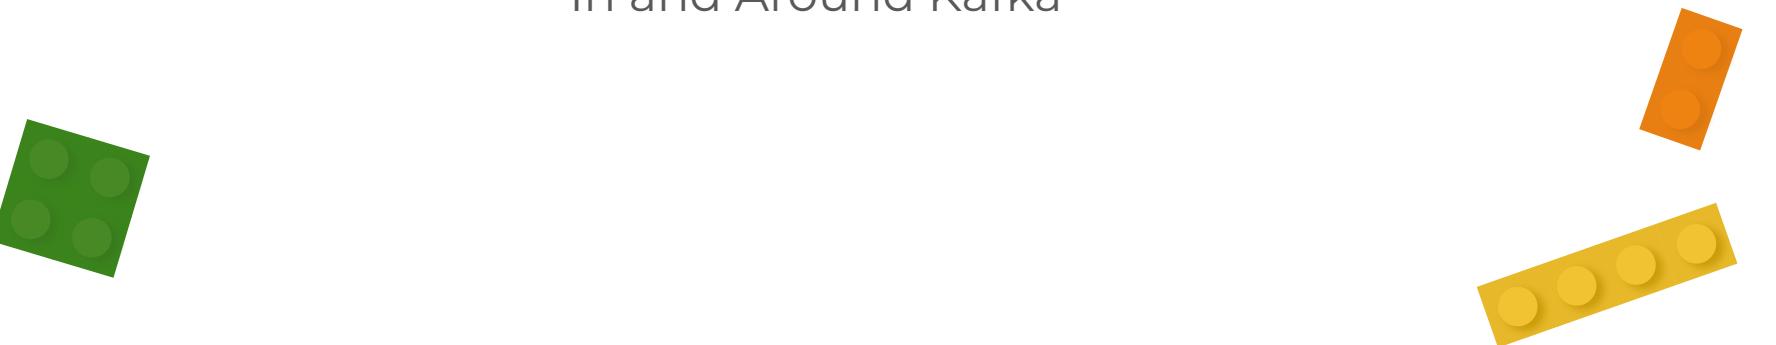
- A framework for integrating data
- 100+ data sources and sinks
- Low- to no-code option
- Built on consumers and producers





Stream Processing

In and Around Kafka





Processing within Kafka

- Consumer/Producer API
 - Lowest level
 - Consume – Process – Produce
 - Manually define state/fault-tolerance
- Kafka Streams
 - Java library for stream processing
 - Built-in state handling and failover
- ksqlDB
 - SQL syntax
 - Kafka Streams under the hood
 - Cloud based offering

`subscribe(), poll(), send(),
flush(), beginTransaction(), ...`

`KStream, KTable,
filter(), map(), flatMap(),
join(), aggregate(), ...`

`CREATE STREAM, CREATE TABLE,
SELECT, JOIN, GROUP BY, SUM, ...`



Flink



Materialize



bytewax



Benthos





Now what?

Use Cases and Inspiration



dfine@confluent.io

@TheDanicaFine

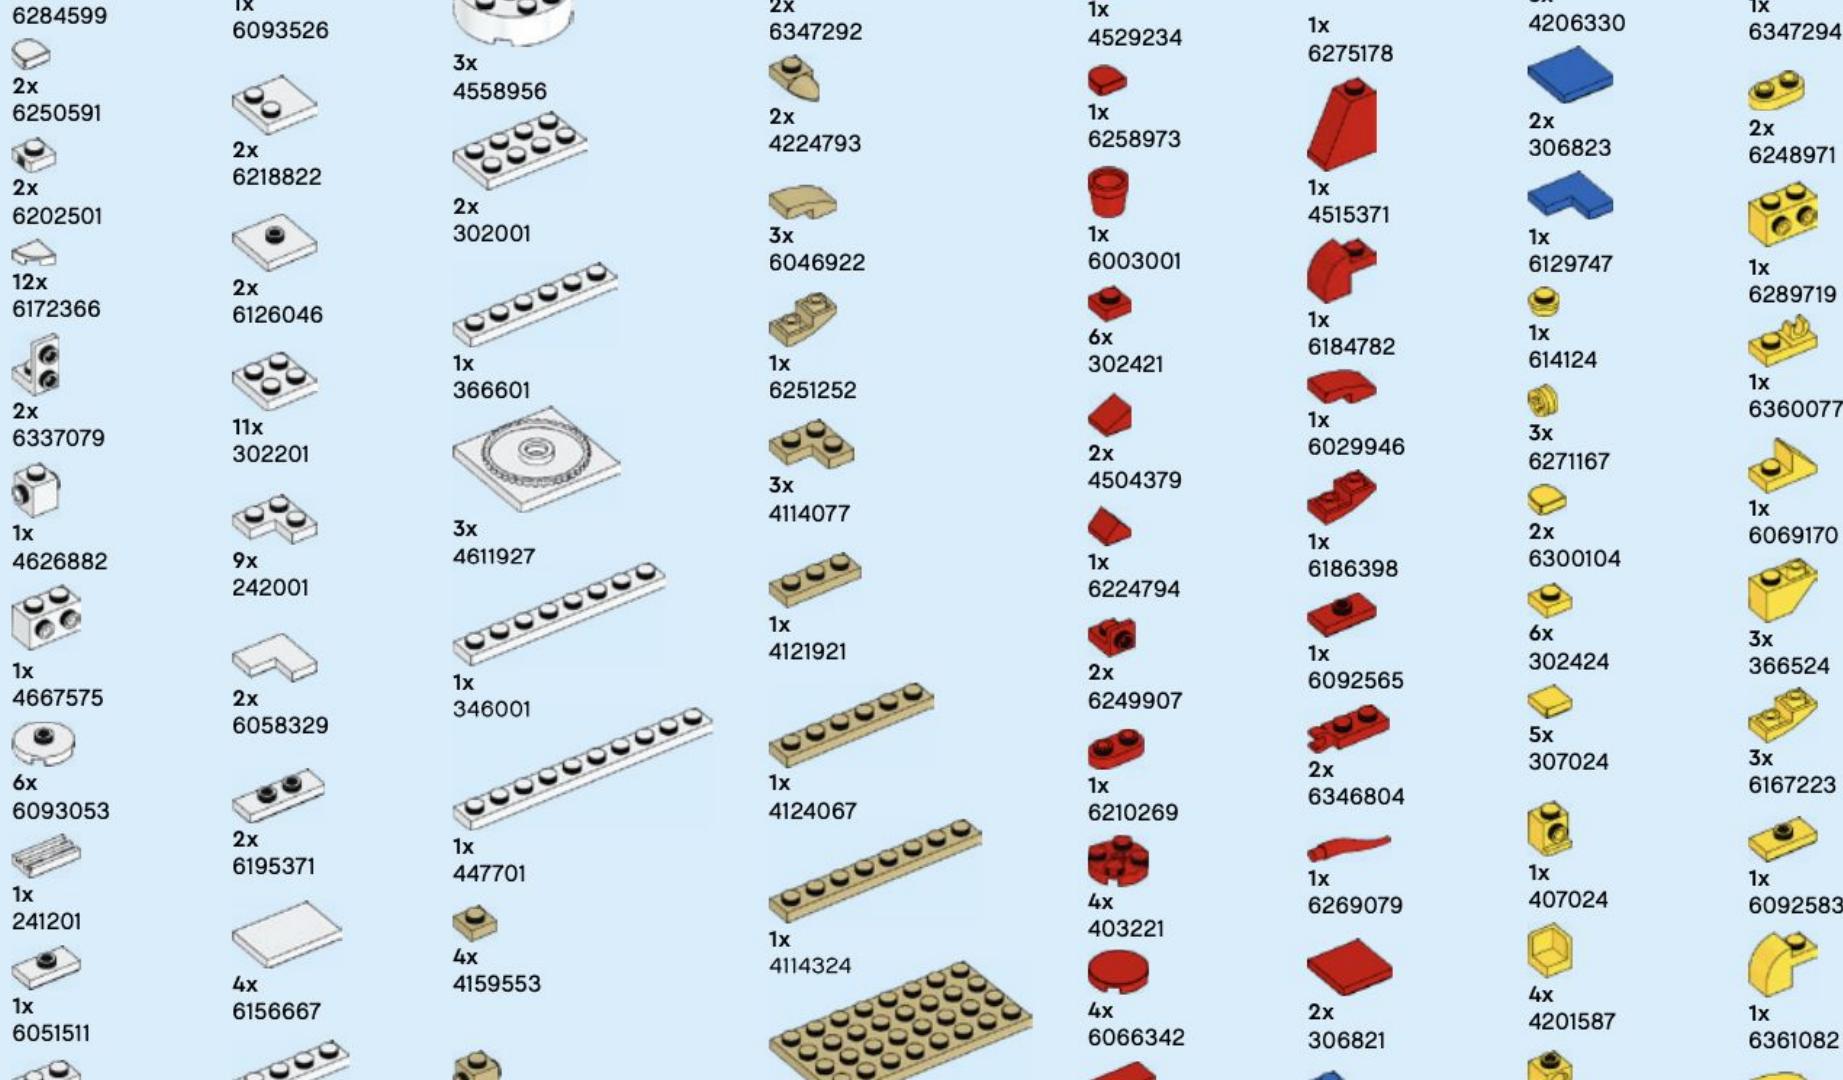
linkedin.com/in/danica-fine/



dfine@confluent.io

@TheDanicaFine

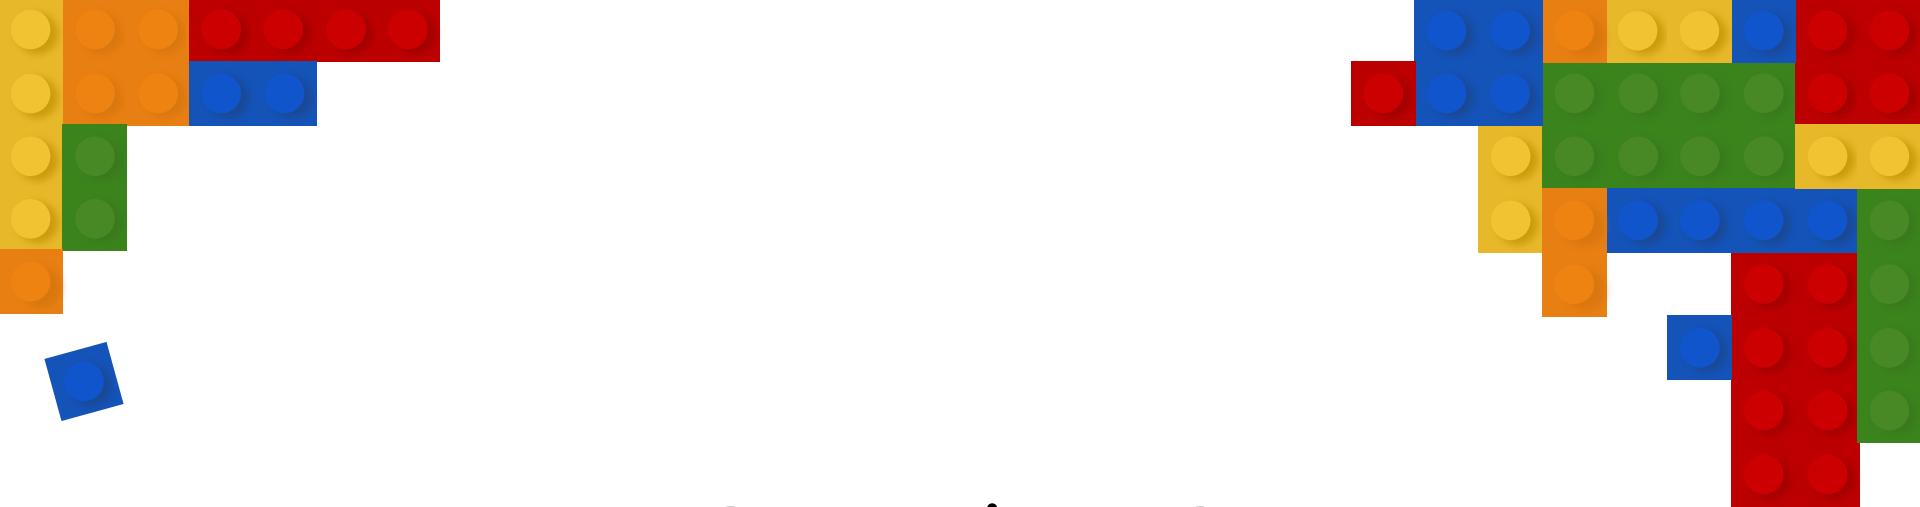
linkedin.com/in/danica-fine/



What will YOU build?

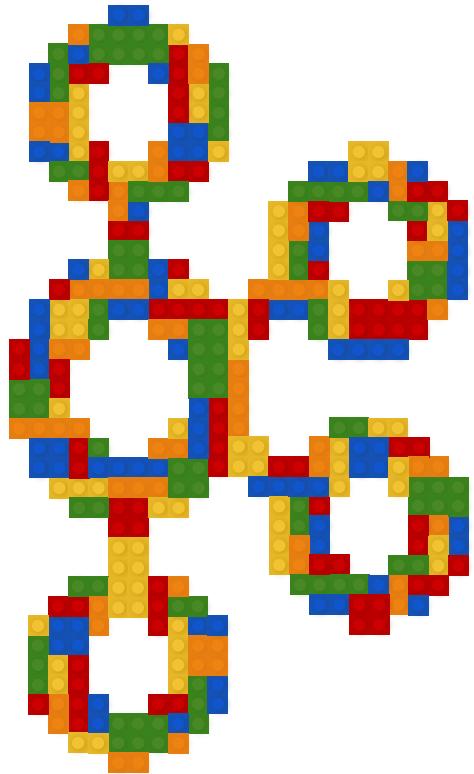


LinkTree Resources



Questions?

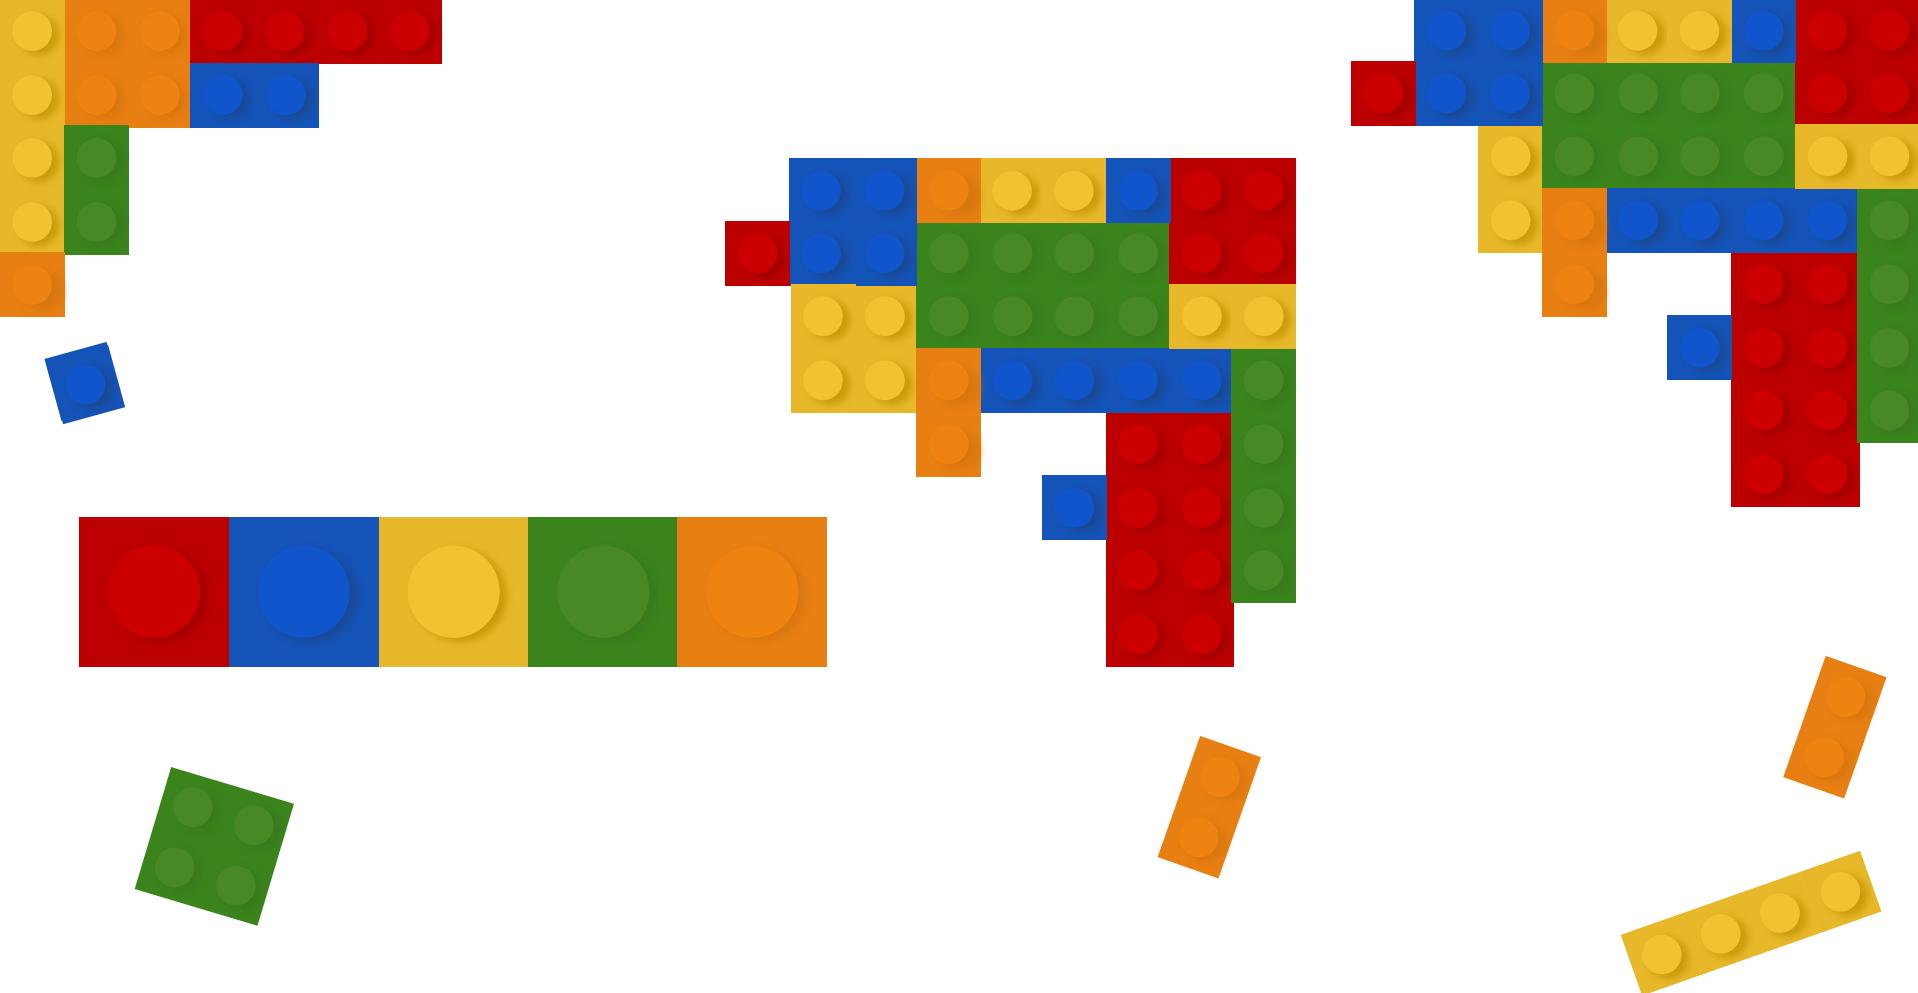




dfine@confluent.io

@TheDanicaFine

linkedin.com/in/danica-fine/



dfine@confluent.io

@TheDanicaFine

linkedin.com/in/danica-fine/

Replication

